

# Homework 1: Numerical Stability and Linear Systems

Due September 21, 2023 (100 points + 20 points extra credit)

**Problem 1** (20 points). Suppose  $\mathbf{x} \in \mathbb{R}^n$  has elements  $x_1, \dots, x_n \in \mathbb{R}$ . The *log-sum-exp* function, common in machine learning code, can be written:

$$f(\mathbf{x}) = \log \sum_{i=1}^n e^{x_i}.$$

- (a) Suppose we implement the formula above numerically. Give examples of cases where the computation of  $f(\mathbf{x})$  can lead to underflow and overflow.
- (b) For  $a \in \mathbb{R}$ , show  $f(\mathbf{x}) = a + \log \sum_{i=1}^n e^{x_i - a}$ . Explain how taking  $a = \max_i x_i$  avoids the issues you identified in the previous part.
- (c) For  $\mathbf{g} = \nabla f(\mathbf{x})$ , show  $g_j = \exp(x_j - \log \sum_{i=1}^n e^{x_i})$ .
- (d) Show that  $\max_i x_i \leq \frac{1}{t} f(t\mathbf{x}) \leq \max_i x_i + \frac{1}{t} \log n$  for any  $t > 0$ . How would you use  $f(\cdot)$  as a differentiable approximation of the map  $\mathbf{x} \mapsto \max_i x_i$ ?

**Problem 2** (20 points). We can use a *condition number* to estimate how difficult it will be to approximate  $f(\mathbf{x})$  for some  $\mathbf{x}$ . Define the relative condition number of  $f(\cdot)$  as

$$\text{cond}(f, \mathbf{x}) = \lim_{\varepsilon \rightarrow 0} \sup_{\|\Delta \mathbf{x}\| \leq \varepsilon \|\mathbf{x}\|} \frac{|f(\mathbf{x} + \Delta \mathbf{x}) - f(\mathbf{x})|}{\varepsilon |f(\mathbf{x})|}.$$

- (a) Suppose we use the  $\infty$ -norm:  $\|\mathbf{x}\|_\infty = \max_k |x_k|$ . For log-sum-exp function  $f$  from Problem 1, show

$$\text{cond}_\infty(f, \mathbf{x}) = \frac{\max_k |x_k|}{|\log \sum_k e^{x_k}|}.$$

*Hint:* You may wish to use the Taylor expansion of  $f$  about  $\mathbf{x}$ .

- (b) Argue that log-sum-exp is ill-conditioned when  $x_i \approx -\log n$  for all  $i$ .
- (c) Argue that log-sum-exp is well-conditioned when  $\max_i x_i = \max_i |x_i|$ .

**Problem 3** (25 points). Suppose we have a multi-class classification problem, wherein we wish to train a classifier to map from points  $\mathbf{x} \in \mathbb{R}^d$  to one of  $k$  classes. A simple model for this task is to learn a *softmax classifier*, parameterized by an unknown matrix  $A \in \mathbb{R}^{k \times d}$ , which maps points  $\mathbf{x} \in \mathbb{R}^d$  to probability vectors via:

$$\mathbf{p}(\mathbf{x}; A) = \frac{e^{A\mathbf{x}}}{\mathbf{1}^\top e^{A\mathbf{x}}}.$$

Here, we take the convention that  $e^{\mathbf{v}}$  is the vector whose elements are  $e^{v_i}$ . We can think of  $\mathbf{p}(\mathbf{x}; A)$  as a vector of *probabilities*, where  $p_i(\mathbf{x}; A)$  is the likelihood that  $\mathbf{x}$  is in class  $i$ .

Suppose we are given a dataset  $\{(\mathbf{x}_i, c_i)\}_{i=1}^N$ , where  $c_k \in \{1, \dots, k\}$ . The *negative log likelihood* loss associated to  $A$  is:

$$\ell(A) = -\frac{1}{N} \sum_{i=1}^N \log p_{c_i}(\mathbf{x}_i; A).$$

Notice  $\ell$  is *small* when our classifier correctly classifies all the points in the dataset.

- (a) Take  $X \in \mathbb{R}^{d \times N}$  to be the matrix whose columns are the  $\mathbf{x}_i$ 's, and suppose  $C \in \{0, 1\}^{k \times N}$  satisfies

$$C_{ij} = \begin{cases} 1 & \text{if } i = c_j \\ 0 & \text{otherwise.} \end{cases}$$

Show  $\ell(A) = \frac{1}{N}(\log \mathbf{1}^\top e^{AX})\mathbf{1} - \frac{1}{N}\text{tr}(C^\top AX)$ , where  $\mathbf{1}$  denotes the vector of all ones and matrix exponentiation is *elementwise*. What is the gradient of  $\ell$  with respect to  $A$ ?

- (b) Implement  $\ell(A)$  and  $\nabla_A \ell(A)$  in `nllSoftmax` in `hw1.py`. Make sure your implementation is stable to large values in  $A\mathbf{x}$ . The assignment includes a dataset and gradient descent code to test the resulting model.

**Problem 4** (10 points). Provide an algorithm for computing the determinant of a matrix  $A \in \mathbb{R}^{n \times n}$  in  $O(n^3)$  time. Justify its correctness and that it works for *all possible* square matrix inputs  $A$ , up to numerical precision.

**Problem 5** (25 points). In the *regression* problem, we are given a number of pairs  $\{(\mathbf{x}_i, y_i)\}_{i=1}^N \subset \mathbb{R}^d \times \mathbb{R}$  and wish to find a function  $f(\cdot)$  so that  $f(\mathbf{x}_i) \approx y_i$ .

- (a) Suppose we seek to fit a function of the form  $f(\mathbf{x}) = \mathbf{a}^\top \mathbf{x}$ . Provide a  $d \times d$  linear system of equations for recovering  $\mathbf{a} \in \mathbb{R}^d$  via least-squares with Tikhonov regularization. Your system should be in terms of  $X \in \mathbb{R}^{d \times N}$ , the matrix whose columns are the  $\mathbf{x}_i$ 's, as well as  $\mathbf{y} \in \mathbb{R}^N$ , the vector whose elements are the  $y_i$ 's.
- (b) Show that  $\mathbf{a} \in \mathbb{R}^d$  from the previous part can be written  $X\mathbf{c}$  for some  $\mathbf{c} \in \mathbb{R}^N$ .
- (c) Give a  $N \times N$  linear system of equations to recover  $\mathbf{c}$ . You can assume  $X^\top X$  is invertible.
- (d) Now, suppose we instead use the form  $f_\phi(\mathbf{x}) = \mathbf{a} \cdot \phi(\mathbf{x})$ , for some function  $\phi : \mathbb{R}^d \rightarrow \mathbb{R}^k$  and an unknown  $\mathbf{a} \in \mathbb{R}^k$ . Define  $K(\mathbf{x}, \mathbf{y}) = \phi(\mathbf{x}) \cdot \phi(\mathbf{y})$ . Give an algorithm for evaluating the least-squares fit  $f_\phi(\mathbf{z})$  at an arbitrary  $\mathbf{z} \in \mathbb{R}^d$  using only evaluations of  $K$ —*not* evaluations of  $\phi$ —as well as linear algebra operations.
- (e) Which matrix factorization would you use for the linear system of equations in part (d)?

**Problem 6** (20 points, extra credit). A typical model of floating-point arithmetic assumes *multiplicative error*. Suppose we attempt to compute  $c = a \text{ op } b$  where  $\text{op} \in \{+, -, \times, \div\}$ . We will assume after rounding that the computer provides us with a value  $\hat{c}$  satisfying  $\hat{c} = c \cdot (1 + \delta)$ , where  $|\delta| < u$  for some fixed constant  $u$ . Applying multiple arithmetic operations incurs *different*  $\delta$ 's:

$$\widehat{(a + b) + c} = ((a + b)(1 + \delta_1) + c)(1 + \delta_2).$$

- (a) Suppose we have a list of nonnegative values  $w_1, \dots, w_n \geq 0$  and wish to compute  $s_k = \sum_{i=1}^k w_i$ ; our code uses a `for` loop, using the relationship  $s_i = s_{i-1} + w_i$  with  $s_1 = w_1$ . Prove the bound  $|\hat{s}_{n-1} - s_{n-1}| \leq u \sum_{i=1}^{n-1} (n - i)w_i + O(u^2)$ .

*Hint:* First show  $|\hat{s}_i - s_i| \leq u \sum_{j=1}^i s_j + O(u^2)$  by induction.

- (b) For some  $a \in \mathbb{R}$ , define  $w_i = e^{x_i - a}$ . Assuming exponentiation also generates multiplicative error, justify the bound  $|\hat{w}_i - w_i| \leq ((1 + |a - x_i|)u + O(u^2))w_i$ .

*Note:* For all parts of this problem, you can assume  $\hat{w}_i \geq 0$  for all  $i$ .

- (c) Suppose we compute  $p_k = \sum_{i=1}^k e^{x_i - a}$  via the recurrence  $p_i = p_{i-1} + w_i$  and  $w_i = e^{x_i - a}$  (with  $p_0 = 0$ ). Accounting for rounding error, using the previous two parts, justify the inequality

$$|\hat{p}_{n-1} - p_{n-1}| \leq u \sum_{i=1}^{n-1} (2 + |a - x_i| + n - i)w_i + O(u^2).$$

- (d) For simplicity, suppose we reshuffle so that  $x_n = \max_i x_i$ . Then, Problem 1 computes

$$f(\mathbf{x}) = x_n + \log \left( 1 + \sum_{i=1}^{n-1} e^{x_i - x_n} \right).$$

Suppose we code up this formula to approximate  $z = x_n + \log(1 + p_{n-1})$ . If we use Python's `log1p` function, we incur error once:  $\log(1 + x) = (\log(1 + x))(1 + \delta)$ . Derive the bound

$$|\hat{z} - z| \leq u [2|z| + n - x_{\min}] + O(u^2).$$

- (e) If we repeat the analysis above without the log-sum-exp trick from Problem 1, we will find a different inequality  $|\hat{z} - z| \leq u [|z| + n + 1] + O(u^2)$ . You've done enough annoying computations: We will not ask you to derive this bound! Based on this formula and the solution to the previous part, explain how shifting from Problem 1 does *not* yield an obvious improvement to accuracy.